



**TECH**  
ELEVATOR

Java

# Intro to Databases

## Ordering, Limiting, and Grouping

Module 2 - Week 1

- Welcome to Module 2
  - Databases
  - Don't forget about Java, DAOs and JDBC coming up.
  - Java Katas added to Module 1
  - Think about side projects (collaborations are okay!).



Week	Topics
Week 0	Module 2 Orientation / PostgreSQL
Week 1	Intro to databases / Ordering, limiting, and grouping
Week 2	SQL joins / Insert, update, and delete / Database design
Week 3	Data Access / Data security
Week 4	DAO testing
Week 5	<b>Mid-module project</b>
Week 6	Postman / NPM / Networking and HTTP / Consuming RESTful APIs
Week 7	Server-side APIs
Week 8	Securing APIs
Week 9	<b>End-of-module project</b>
Week 10	<b>Assessment</b>

# SQL and NoSQL Databases

## Pros of Relational Databases

- Great for structured data
- Use of an existing query language (SQL)
- Great for complex queries
- Easy data navigation
- High level of data integration, due to relationships and constraints among tables
- Transactions are secure

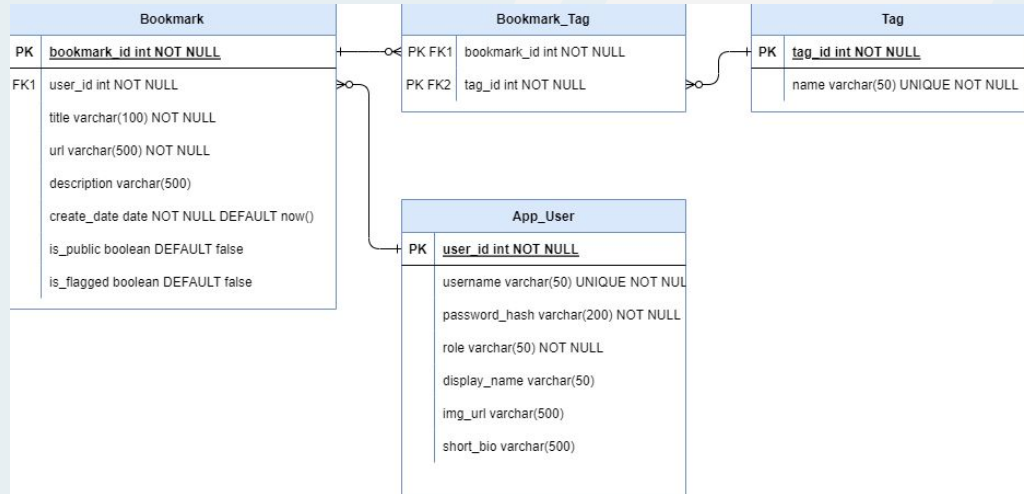
## Pros of Non-Relational Databases

- Flexible data model
- Rapid adaptation to changing requirements: dynamic changes to a item do not affect the other items
- Storage of huge amount of data with little structure
- High performance

[Source: towardsdatascience.com](https://towardsdatascience.com)

# Bookmark Manager Sample Application

- Log in as a user
- View a list of public bookmarks
- Review any flagged bookmarks
- Manage a list of tags that users associate with a bookmark
- Review ERD
  - users can have many bookmarks that they're associated with
  - users may assign a bookmark many different tags
  - users may use the same tag for many bookmarks



# (R)DBMS

A Relational Database Management System ( (R)DBMS) is a software application designed to manage a database. It has four basic functions

1. Data Definition
2. Data Storage
3. Data Retrieval
4. Administration

RDBMSs include databases like Oracle, Microsoft SQL Server, PostgreSQL, MySQL, are relational, and are commonly called **SQL Databases**.

**NoSQL** Databases are those that do not use a relational structure, instead they structure data specific to the problem they are designed to solve. NoSQL databases include MongoDB, Cassandra, Google BigTable, HBase, DynamoDB, and Firebase.

[DB-Engines](#) has a ranking measuring popularity of current DBMS platforms.

# Structured Query Language (SQL)

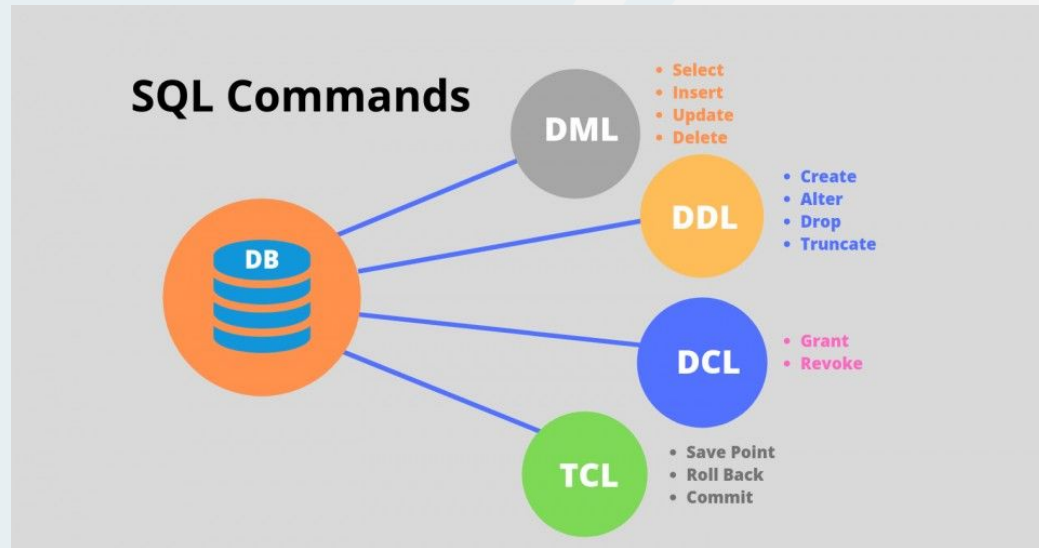
SQL is a declarative programming language used to manage a database and its data.

A **declarative programming language** specifies what actions should be performed rather than how to perform those actions.

SQL Consists of 3 sub-languages

1. **DDL - Data Definition Language** - defines the structure of the data
2. **DML - Data Manipulation Language** - query and modify the data
3. **DCL - Data Control Language** - used to administer the database

- Data definition language (DDL): modify or alter the structure of the database.
- Data manipulation language (DML): perform operations like storing data in database tables, modifying and deleting existing rows, retrieving data, or updating data.
- Data control language (DCL): perform operations like giving or removing database access for a user.
- Transaction Control Language (TCL): control the execution of a transaction.





**SQL** - Structured Query Language : a language that lets you access and manipulate databases

**ANSI-SQL** - A standard that databases must follow to be considered a SQL Database.

All SQL databases support the ANSI-SQL language, however, most databases extend it with their own proprietary additions.

An **entity** is a set of data being stored a *table*.

A **Table** defines a set of data elements and the structure to store them. Tables are structured into *columns* and *rows*.

**Columns** - attributes of a table and define the name and data type. A table has a set number of defined columns.

**Rows** - the data being stored. A table has an unlimited number or rows.

A **Cell** is the location where a *column* and *row* intersect, and is used to refer to a specific value or row of data (*entity*).

# Column

id	name	countrycode	district	population
3793	New York	USA	New York	8008278
3794	Los Angeles	USA	California	3694820
3795	Chicago	USA	Illinois	2896016
3796	Houston	USA	Texas	1953631
3797	Philadelphia	USA	Pennsylvania	1517550
3798	Phoenix	USA	Arizona	1321045
3799	San Diego	USA	California	1223400
3800	Dallas	USA	Texas	1188580
3801	San Antonio	USA	Texas	1144646
3802	Detroit	USA	Michigan	951270

## Column

- Tables have a set number.
- Define the data the table will hold
- Provides a label for each part of the data being stored

## Columns on this Table

id, name, countrycode, district, population

# Row

id	name	countrycode	district	population
3793	New York	USA	New York	8008278
3794	Los Angeles	USA	California	3694820
3795	Chicago	USA	Illinois	2896016
3796	Houston	USA	Texas	1953631
3797	Philadelphia	USA	Pennsylvania	1517550
3798	Phoenix	USA	Arizona	1321045
3799	San Diego	USA	California	1223400
3800	Dallas	USA	Texas	1188580
3801	San Antonio	USA	Texas	1144646
3802	Detroit	USA	Michigan	951270

## Row

- Tables have a unlimited number (0...n)
- Contain the data
- Has a value for each column

# Cell

id	name	countrycode	district	population
3793	New York	USA	New York	8008278
3794	Los Angeles	USA	California	3694820
3795	Chicago	USA	Illinois	2896016
3796	Houston	USA	Texas	1953631
3797	Philadelphia	USA	Pennsylvania	1517550
3798	Phoenix	USA	Arizona	1321045
3799	San Diego	USA	California	1223400
3800	Dallas	USA	Texas	1188580
3801	San Antonio	USA	Texas	1144646
3802	Detroit	USA	Michigan	951270

## Cell

- The intersection of a column and row
- Used to identify a specific row of data

In this example we would identify the row we want to access by saying the ROW where the CELL in the COLUMN labelled 'name' has the value 'Chicago'

- Character Types
  - `char(#)` - character. # defined the length of the data.
  - `varchar(#)` - varying character. # defined the length of the data.
  - `text` - text based data that is not limited by a predefined size
- Numeric Types
  - `int` - similar to Java's int
  - `serial` - works similar to the integers except these are automatically generated in the columns by PostgreSQL
  - `numeric(d,p)` - floating point numbers with d number of digits and p number of decimal places
- Other types
  - `boolean` - true/false
  - `date` - yyyy-mm-dd
  - `time` - hh:mm:ss
  - `timestamp` - yyyy-mm-dd hh:mm:ss

The most basic SQL statement is a `SELECT` query, and it follows the following format:

```
SELECT [column], [column-n] FROM [table];
```

- `SELECT [column], [column-n]` indicates which columns that you want returned from your query.
- `FROM [table]` indicates the table from which to retrieve the data.

```
SELECT city_name, population FROM city;
```

```
SELECT * FROM city;
```

- `SELECT *` indicates ALL columns returned from your query.

city  
table

city_id	city_name	state_abbreviation	population
3793	New York	NY	8008278
3794	Los Angeles	CA	3694820
3795	Chicago	IL	2896016
3796	Houston	TX	1953631
3797	Philadelphia	PA	1517550
3798	Phoenix	AZ	1321045
3799	San Diego	CA	1223400

The **WHERE** clause is used to filter the result set based on criteria rules. Can include:

- **=, <>, !=, >, >=, <, <=**
- **IS NULL, IS NOT NULL**
- **IN(values), NOT IN(values)**
- **LIKE** (with wildcard character)
- **BETWEEN** value **AND** value

```
SELECT city_name, population FROM city WHERE state_abbreviation = 'PA';
```

```
SELECT city_name, population FROM city WHERE population > 3000000
```

city  
table

city_id	city_name	state_abbreviation	population
3793	New York	NY	8008278
3794	Los Angeles	CA	3694820
3795	Chicago	IL	2896016
3796	Houston	TX	1953631
3797	Philadelphia	PA	1517550
3798	Phoenix	AZ	1321045
3899	Pittsburgh	PA	301286



- Write and execute simple select statements using SELECT, FROM, and WHERE
- **Use aggregate information using GROUP BY to group rows together**
- Use the ORDER BY clause to order results from the database
- Write select statements using subqueries

SQL is a declarative language and does not run from top to bottom and left to right. The order it runs is:

1. `FROM` clause - The database needs to know which table(s) you're selecting from first of all
2. `WHERE` clause - The database then needs to know which rows you'll work with
3. `GROUP BY` clause - The database then groups those rows according to your `GROUP BY` clause
4. `SELECT` clause - The database then collapses those rows down and selects the columns that you want data from
5. `ORDER BY` clause - The database orders the rows in the order that you ask for
6. `LIMIT` clause - The database only returns the number of resulting rows that you want

**Select**

**From**

**Where**

**Group by**

**Having**

**Order by**

**Limit**

**Some**

**French**

**Waiters**

**Grow**

**Healthy**

**Oranges &**

**Lemons**

- AS can be used with a column name to give it an alias (new name)

```
SELECT city_name AS name_of_city FROM city;
```

Data Output	Explain	Mes
name_of_city character varying (50)		
1	Abilene	
2	Akron	

```
SELECT city_name || ' ' || state_abbreviation AS city_and_state FROM city;
```

Data Output	Explain	Mess
city_and_state text		
1	Abilene TX	
2	Akron OH	

- DISTINCT can be used with a column name to return only unique values from that column.

```
SELECT city_name from city;
```

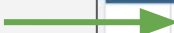
Data Output	Explain	Messa
city_name character varying (50)		
295	Springfield	
296	Springfield	
297	Springfield	
298	St. Louis	

```
SELECT DISTINCT city_name from city;
```

Data Output	Explain	Messa
city_name character varying (50)		
7	St. Petersburg	
8	Springfield	
9	St. Louis	

- In Postgres, we can concatenate the values across multiple columns into a single field using the || operator to concatenate strings.
- We can use the AS keyword to give the concatenated column a name.

```
SELECT city_name || ', ' || state_abbreviation  
AS city_state FROM city  
WHERE state_abbreviation IN ('PA', 'OH', 'NY');
```



	Data Output	Explain	Me
	<b>city_state</b> text		
1	Akron, OH		
2	Albany, NY		
3	Allentown, PA		
4	Buffalo, NY		
5	Cincinnati, OH		
6	Cleveland, OH		
7	Columbus, OH		
8	Dayton, OH		
9	Harrisburg, PA		
10	New York City, NY		

Aggregate function performs a calculation on a set of values and returns a single value.

- AVG returns the average value of a numeric column
- SUM returns the total sum of a numeric column
- COUNT returns the number of rows matching criteria
- MIN returns the smallest value of the selected column
- MAX returns the largest value of the selected column

Except for `COUNT (*)`, aggregate functions ignore null values. It's important to be aware of the difference between `COUNT (*)` and `COUNT` on a specific field.

If you specify a column name to count, only the rows in the table that have a value for that column will be returned. For instance, in our state data, this query returns 51 rows, while changing it to use `COUNT (*)` instead returns 56. This is because only 51 rows have a value (that do not have a `NULL` value) for `state_nickname`.

```
SELECT COUNT(state_nickname) FROM state;
```

51

```
SELECT COUNT(*) FROM state;
```


56

An aggregate function performs a calculation on one or more values and returns a single value.

Very often with aggregate functions such as `SUM`, we want the results to be grouped by the value of some attribute. The `GROUP BY` clause groups the result set into groups of values and the aggregate function returns a single value for each group.

For example if our database contains city information, including state abbreviation and population, we can report the total population in each state abbreviation by using `SUM` and `GROUP BY` the state abbreviation .

```
SELECT state_abbreviation, SUM(population) as sum_city_population
FROM city
GROUP BY state_abbreviation;
```





	Data Output	Explain	Messages	Notif
	state_abbreviation character (2)		sum_city_population bigint	
1	AK		319276	
2	AL		898351	
3	AR		197312	
4	AS		3656	
5	AZ		4286236	
6	CA		20078730	
7	CO		2726421	
8	CT		633960	
9	DC		705740	

city  
table

city_name	state_abbreviation	population
New York	NY	8008278
Allentown	PA	121442
Syracuse	NY	142327
Cleveland	OH	381009
Philadelphia	PA	1517550
Columbus	OH	898553
Pittsburgh	PA	301286

```
SELECT
state_abbreviation,
MAX(population) as max_city_population
FROM city
GROUP BY state_abbreviation;
```

	state_abbreviation character (2)	max_city_population integer
1	OH	898553
2	NY	8008278
3	PA	1517550





- Write and execute simple select statements using SELECT, FROM, and WHERE
- Use aggregate information using GROUP BY to group rows together
- **Use the ORDER BY clause to order results from the database**
- Write select statements using subqueries

A result set can be sorted using the ORDER BY clause:

```
SELECT col1, col2
FROM tablename
WHERE col1 = 'value'
ORDER BY col1 [ASC | DESC], col2 [ASC | DESC]
```

- Sort columns must exist in the table being queried or can be aliased columns
- Multiple column names can be provided which assigns a priority sort.
- Each column in ORDER BY clause can be specified as ascending (ASC) or descending (DESC). If not specified the default is ASC.

```
SELECT census_region, state_name, population
FROM state ORDER BY census_region, population DESC;
```

By using `LIMIT n` at the end of a query, we can limit the size of our result set to `n` results.

```
SELECT state_name FROM state LIMIT 10;
```

This tends to work best with `ORDER BY` as it allows you to construct lists like the top 10 states by population.

```
SELECT state_name, population  
FROM state ORDER BY population DESC  
LIMIT 10;
```

- Write and execute simple select statements using SELECT, FROM, and WHERE
- Use aggregate information using GROUP BY to group rows together
- Use the ORDER BY clause to order results from the database
- **Write select statements using subqueries**

A subquery is referred to as an inner query and can provide the results of one query as input to another.

- Often used in the WHERE clause

```
SELECT city_name, state_abbreviation FROM city
WHERE state_abbreviation IN (SELECT state_abbreviation FROM state WHERE census_region = 'Northeast');
```

- Can be used in the SELECT clause (more on joins next week)

```
SELECT s.state_abbreviation, (SELECT c.city_name FROM city AS c WHERE c.city_id = s.capital)
FROM state AS s;
```

- Or even in the FROM clause.

```
SELECT s.city_state, s.population
FROM (SELECT city_name || ', ' || state_abbreviation AS city_state, population FROM city) AS s;
```